




A Convolutional Neural Network for Image Classification of Cats and Dogs

Final presentation



INTRODUCTION

Structure



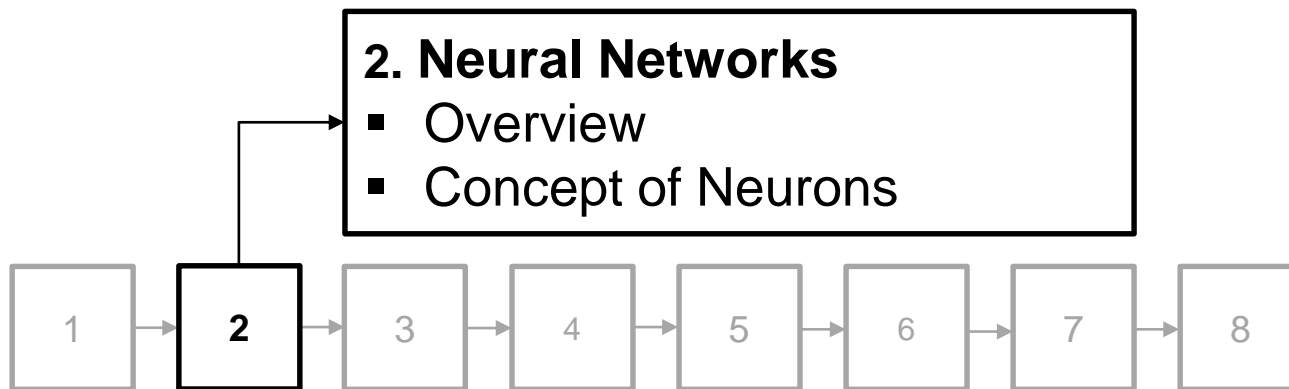
Structure

1. Introduction
2. Neural Networks (NN)
3. Math behind NN
4. Convolutional NN (CNN)
5. Problem
6. Design
7. Evaluation
8. Summary




NEURAL NETWORKS

Structure




Overview



<http://cs231n.github.io/convolutional-networks/>

Concept of Neurons

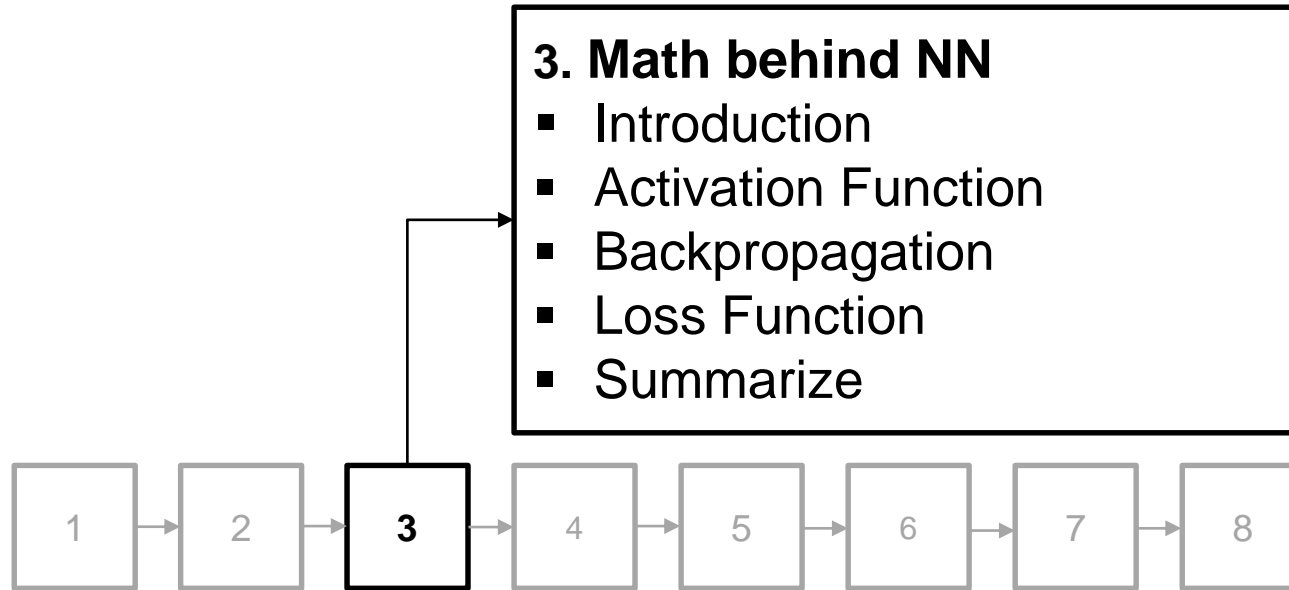


<http://cs231n.github.io/convolutional-networks/>



MATH BEHIND NEURAL NETS

Structure





Introduction

- Actual output, weights
- Activation function
- Measure how much we missed (cost function)
- Multiply error by the Sigmoid slope
- Update weights (backpropagation)

Activation Functions


- Non-Linear
 - Ex: Sigmoid, tanh
- Continuous but not everywhere differentiable function
 - Cons: descent gradient cannot be obtained
 - Ex: Relu



<http://cs231n.github.io/convolutional-networks/>

Back Propagation

- $f(x, y, z) = (x + y)z$
- $x = -2, y = 5, z = -4$
- $q = x + y, \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$
- $f = qz, \frac{\partial f}{\partial q} = z, \frac{\partial q}{\partial z} = q$
- Desired
 - $\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$
- Similarly propagate
 - $\frac{\partial f}{\partial y}$ and $\frac{\partial f}{\partial z}$



<http://cs231n.github.io/convolutional-networks/>



Loss Function

- Squared Error Measure
- SoftMax

Loss Function - Squared Error Measure

- $Error = \frac{1}{2} (Y_{actual} - Y_{predicted})^2$
- Drawbacks
 - No gradient to get from 0.000...1 to 1.
 - To do so it will take quite longer.
 - Deprives NN of probability information.

Summarize

- Actual output, weights
- Activation function
- Measure how much we missed (cost function)
- Multiply error by the Sigmoid slope
- Update weights (backpropagation)

$$l_0 = T_i, W = rand()$$

$$l_1 = f(X_i \cdot W)$$

$$Err = (l_0, l_1)$$

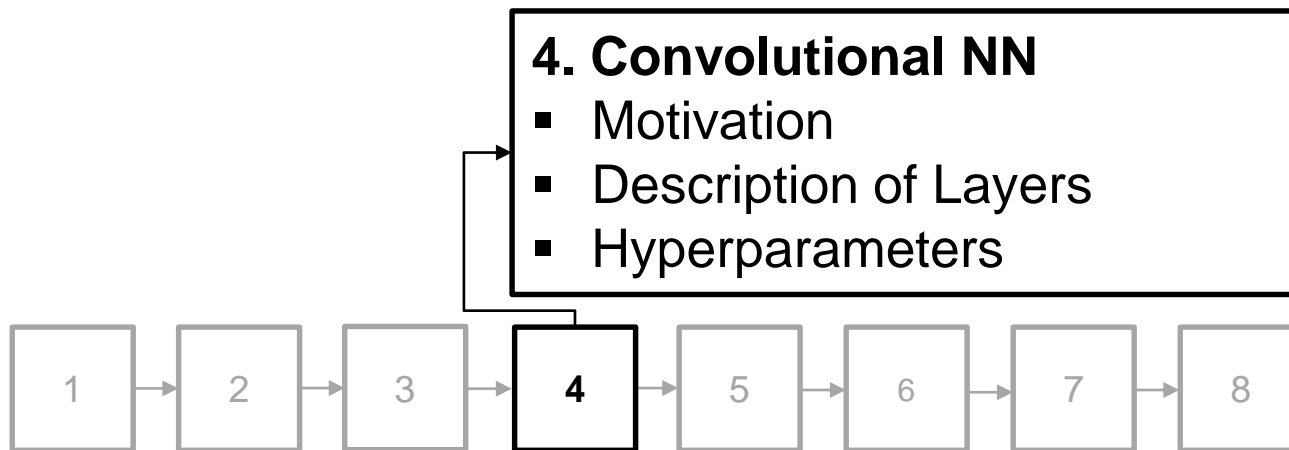
$$\Delta l_1 = Err \times \Delta(f(Err))$$

$$W = W + \alpha(l_0 \cdot \Delta l_1)$$




CONVOLUTIONAL NN

Structure



Motivation


- Number of parameters



Motivation

■ NN

- High number of params




| | | |
|-------|-------|-------|
| x_1 | x_2 | x_3 |
| x_4 | x_5 | x_6 |
| x_7 | x_8 | x_9 |

Number of
weights: 36


■ CNN

- Lower number of params




Number of
weights: 4

Description of Layers



Input Layer

- Image cropping
- Distortions
 - Randomly flipping
 - Randomly changing brightness
 - Randomly changing contrast



Convolutional Layer - Filter

Input

| | | | | | | |
|---|---|---|---|----------|----------|----------|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 2 | 1 | 0 | 0 | 2 | 0 |
| 0 | 2 | 2 | 2 | 0 | 2 | 0 |
| 0 | 1 | 2 | 0 | 2 | 0 | 0 |
| 0 | 2 | 1 | 2 | 1 | 2 | 0 |
| 0 | 1 | 2 | 1 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Filter

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 1 |
| 0 | 2 | 2 |

Bias


| |
|---|
| 1 |
|---|

Output

| |
|---|
| 1 |
|---|



$$\begin{aligned}
 &1*0 + 2*0 + 0*0 + \\
 &0*0 + 0*2 + 0*1 + \\
 &0*0 + 0*2 + 0*2 + \\
 &1 = 1
 \end{aligned}$$

<http://cs231n.github.io/convolutional-networks/>




Convolutional Layer - Parameters

- Input volume size
- Number of filters
- Filter size
- Step size
- Zero padding





Convolutional Layer – Activation function

- Rectified linear
 - Element wise $\max(0, x)$
- Leaky ReLu
 - If $x < 0$ then $f(x) = 0.01x$
 - Non-zero gradient when the input is negative






Pool Layer – Max Pooling

- Reduce the spatial dimension of an image



Norm Layer




- 4D-array
- Normalize each element of this array



Norm Layer

- Normalize each element of this array

- $$a_1 = \left(\left(\frac{R}{\sqrt{R^2 + G^2 + B^2}} \right), \left(\frac{G}{\sqrt{R^2 + G^2 + B^2}} \right), \left(\frac{B}{\sqrt{R^2 + G^2 + B^2}} \right) \right)$$





Local Layer

- Also named fully connected layer

Input


| | | |
|-------|-------|-------|
| x_1 | x_2 | x_3 |
| x_4 | x_5 | x_6 |
| x_7 | x_8 | x_9 |

Input




Softmax-Linear Layer

- Softmax output function
- Cost measure for softmax




Softmax Output Function

- Soft continuous version of Max Function





$$Y_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

- Forces $\sum(Y_i) = 1$.



Softmax Output Function

- $\frac{\delta Y_i}{\delta z_i} = Y_i(1 - Y_i)$
- Nice Simple derivative
- Even though Y_i depends of Z_i
 - Derivative
 - For an individual neuron
 - Of an O/P in respect to I/P is just $Y_i(1 - Y_i)$



Cost Measure for Softmax

- Cross entropy cost function

- $\mathcal{C} = -\sum_j T_j \log Y_j$


- Negative log probability of correct answer

- Maximise the log probability of getting answer right

- Very big gradient when O/P is 1 and target is 0

- $\frac{\delta \mathcal{C}}{\delta Z_i} = T_i - Y_i$

- Slope is -1 when target values and actual value is opposite






Hyperparameters - Learning Rate

- How fast the network trains
- High learning rate
 - Convergence or global minimum finding is problem
- Low learning rate
 - High training times

Hyperparameters - Learning Rate Decay


- Learning rate decay means the learning rate decreases over time
 - higher learning rate is well suited to get close to the global minimum
 - small learning rate is better at fine tuning the global minimum
- Several way
 - Exponential decay, reduction by factor of n
 - Function to decrease the learning rate by 4%

Hyperparameters - Overfitting or Underfitting



Hyperparameters - Weight Penalty



- Adding λ to penalise
 - Keeps weight small
 - Big error derivatives



- $C = E + \frac{\lambda}{2} \sum_{i=1} w_i^2$
- $\frac{\partial C}{\partial w_i} = \frac{\partial E}{\partial w_i} + \lambda w_i$
- When $\frac{\partial C}{\partial w_i} = 0$;
 - $w_i = -\frac{1}{\lambda} \frac{\partial E}{\partial w_i}$
 - At minimum of cost function if $\frac{\partial E}{\partial w_i}$ is large, the weights are large

Hyperparameters - Weight Penalty


- Preventing network from the weights it does not need
 - Don't have a lot of weights not doing anything
 - So output changes more slowly as input changes.
- Putting half the weight on each and not on one





PROBLEM

Structure





TU Clausthal

Introduction



Data

- Images of cats and dogs
- File format is *.jpg
- Color space is RGB




dog1.jpg



cat10.jpg

Data

- 25,000 images
 - 12,500 of dogs
 - 12,500 of cats
- Avg. file size
 - 22.34 kB






DESIGN


Structure

6. Design

- Implementation of CNN architecture
- System model
- Implemented architectures



Implementation of CNN Architecture



Implementation of CNN Architecture

- TensorFlow was developed by Google Brain team
- Version 1.0
- Use cases
 - Handwritten patterns, image recognition, Word2Vec
- Input data
 - Audio, image, text
- Used techniques
 - Linear classifiers, NN, CNN




Implementation of CNN Architecture

- Input: Raw model
- Graph
 - Architecture of nodes and edges (like NN structure)
 - Session is placed on device
 - Initialise variables randomly
 - Run
 - Let tensors pass through the graph
- Output: Trained model



System Model



System Model -Train vs. Test Data

- Split data
 - Train data
 - 20,000 images (80 percent)
 - Divide into 5 batches containing 4,000 each
 - Test data
 - 5,000 images (20 percent)

Process images

- Resize to $32 * 32 * 3 = 3,072$
- Convert to array
 - $25,000 * 3,073$



dog1.jpg



cat10.jpg

Process images

- Resize to $32 * 32 * 3$
- Convert to array
 - $25,000 * 3,073$
- Example
 - 1; 22; 11; 123; ...
 - 0; 256; 255; 0; ...



dog1.jpg



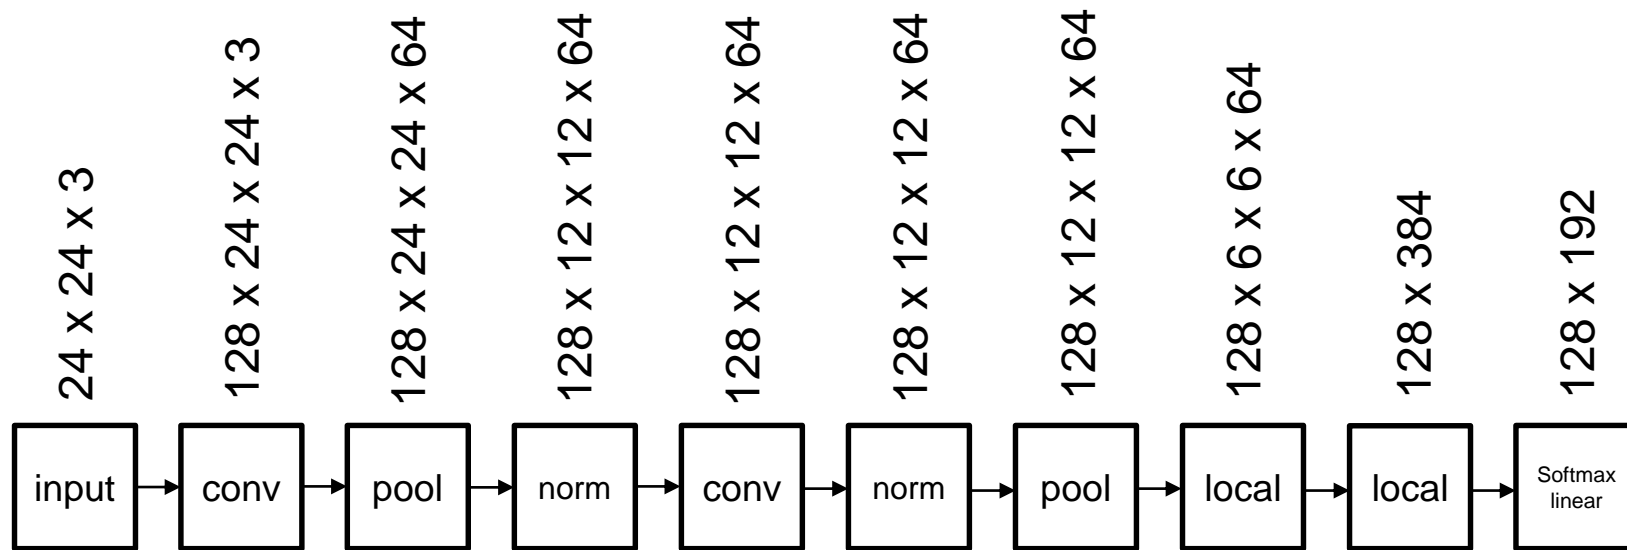
cat10.jpg



System Model - Random distorsion



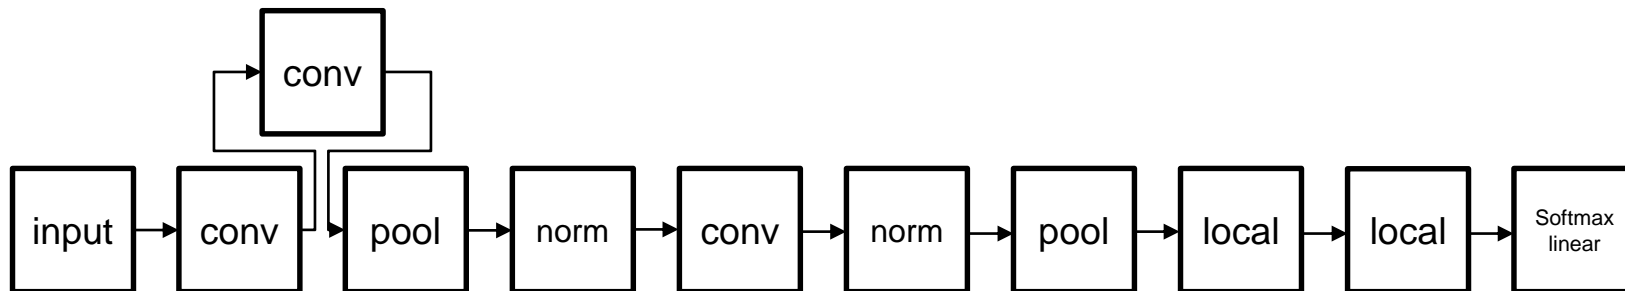
System Model - Structure of CNN



Output: 128×2

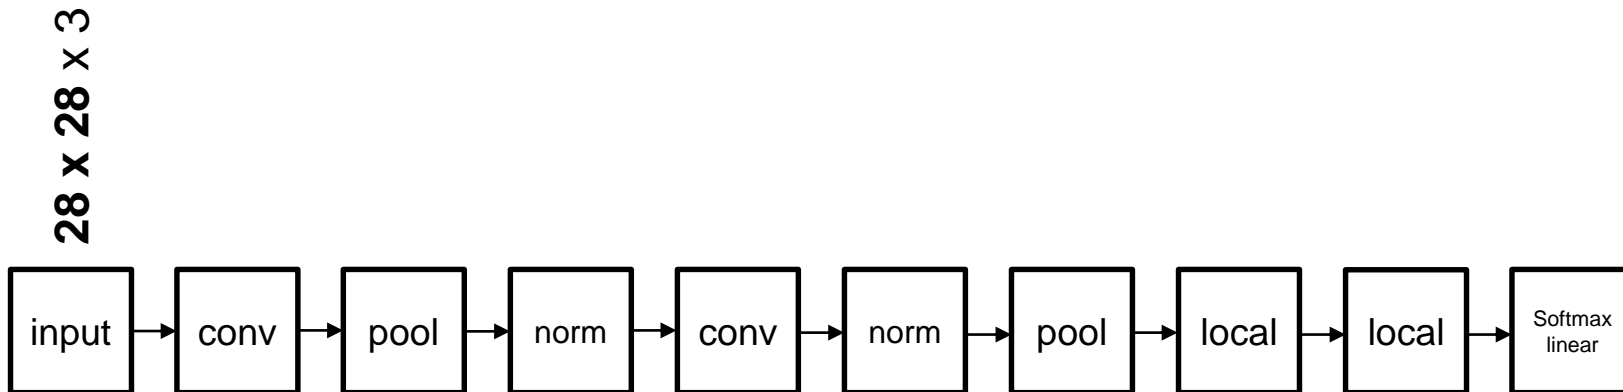
Implemented Architectures – Added Conv Layer

- Input: $128 \times 24 \times 24 \times 3$
- Output: $128 \times 24 \times 24 \times 3$



Implemented Architectures – Increased size

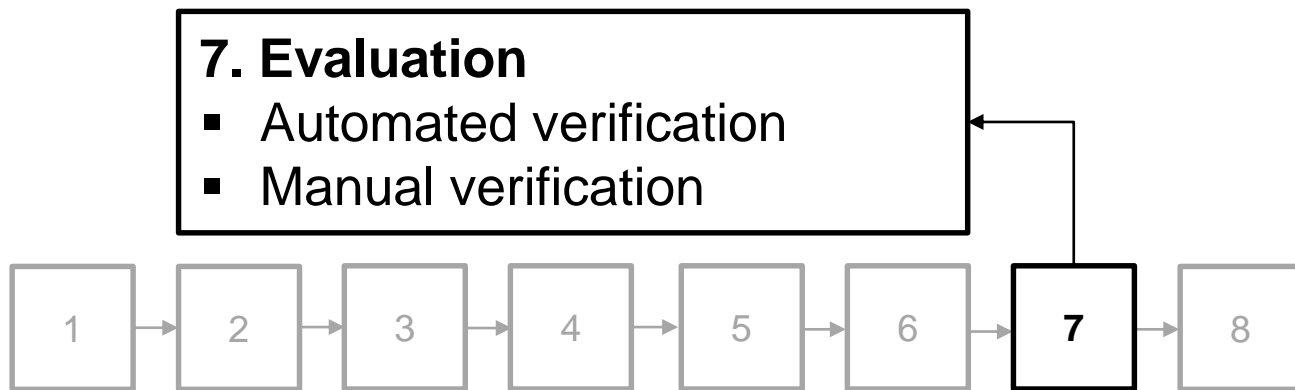
- Original input: $24 \times 24 \times 3$
- New input: $28 \times 28 \times 3$





EVALUATION

Structure




Automated Verification

| | Number of steps | Total loss | Time | Machine |
|-----------------------------|-----------------|------------|------------|---------|
| Standard 100k | 99,900 | 0.1132 | 6h 15m 50s | Windows |
| Standard 40k | 40,600 | 0.3316 | 8h 26m 58s | Linux |
| Additional ConvLayer | 13,500 | 0.3128 | 9h 19m 34s | Linux |
| Increased size | 30,100 | 0.3446 | 8h 37m 30s | Linux |

Automated Verification

- Accuracy of 85 percent




Manual Verification

- For 5000 images accuracy of 97 percent
- Manually verified 100 images
 - Seven were predicted wrong
 - 93 were predicted correctly

Manual Verification – Correctly Predicted



Manual Verification – Wrongly Predicted




Manual Verification – Confusing Images





SUMMARY

Structure



Summary

1. Introduction
2. Neural Networks (NN)
3. Math behind NN
4. Convolutional NN (CNN)
5. Problem
6. Design
7. Evaluation
8. Summary




QUESTIONS

Quellen


- <http://cs231n.github.io/convolutional-networks/>
- https://www.tensorflow.org/tutorials/deep_cnn/
- Maas, Andrew L., Awni Y. Hannun, and Andrew Y. Ng. "Rectifier nonlinearities improve neural network acoustic models." *Proc. ICML*. Vol. 30. No. 1. 2013.




Standard 100k




Standard 40k




Increased Image Size



Added Convolutional Layer




Structure of the CNN we used




Output: 128 x 2


Learning rate



Cross-entropy



Total loss



Total loss after 100k steps roughly above 0.1