



A Convolutional Neural Network for Image Classification of Cats and Dogs

Status update




Structure

- Neural Nets (NN)
- Convolution NN (CNN)
- Hyperparameters
- Problem
- Evaluation




NEURAL NETS

Introduction – Neural Nets



Quelle: <http://cs231n.github.io/convolutional-networks/>

Neuron Model



Mathematical view

- Input, weights
- Compute Sigmoid (activation function)
- Measure how much we missed (cost function)
- Multiply error by the Sigmoid slope
- Update weights (backpropagation)

$$l_0 = X_i, W = \text{rand}()$$

$$l_1 = \text{Sig}(X_i \cdot W)$$

$$\text{Err} = l_0 - l_1$$

$$\Delta l_1 = \text{Err} \cdot \Delta(\text{Sig}(\text{Err}))$$



$$W = W + \alpha(l_0 \cdot \Delta l_1)$$

Activation Functions

- Non-Linear
 - Sigmoid, tanh
- Continuous but not everywhere differentiable function
 - Relu


Activation functions - Sigmoid

- Motivation
 - Not telling in which direction should we move in.
 - Non-differentiability at certain points



Motivation for CNN


- Number of parameters



Motivation for CNN

■ NN

- High number of params




| | | |
|-------|-------|-------|
| x_1 | x_2 | x_3 |
| x_4 | x_5 | x_6 |
| x_7 | x_8 | x_9 |

Number of
weights: 36

■ CNN

- Lower number of params



Number of
weights: 4



CONVOLUTIONAL NN




TensorFlow




- Developed by Google Brain team
- Use cases
 - Handwritten patterns, image recognition, Word2Vec
- Input data
 - Audio, image, text
- Used techniques
 - Linear classifiers, NN, CNN

Structure of the CNN we used



Input layer

- Image cropping
- Distortions
 - Randomly flipping
 - Randomly changing brightness
 - Randomly changing contrast



Convolutional layer - Filter

Input

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 2 | 1 | 0 | 0 | 2 | 0 |
| 0 | 2 | 2 | 2 | 0 | 2 | 0 |
| 0 | 1 | 2 | 0 | 2 | 0 | 0 |
| 0 | 2 | 1 | 2 | 1 | 2 | 0 |
| 0 | 1 | 2 | 1 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Filter

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 1 |
| 0 | 2 | 2 |

Bias


| |
|---|
| 1 |
|---|

Output

| |
|---|
| 1 |
|---|



$$\begin{aligned}
 &1*0 + 2*0 + 0*0 + \\
 &0*0 + 0*2 + 0*1 + \\
 &0*0 + 0*2 + 0*2 + \\
 &1 = 1
 \end{aligned}$$

<http://cs231n.github.io/convolutional-networks/>




Convolutional layer - Parameters

- Input volume size
- Number of filters
- Filter size
- Step size
- Zero padding




Convolutional layer – Activation function

- Rectified linear
 - *Element Wise*: $\max(0, x)$
 - Leaky Relu
 - If $x < 0$, Output = $0.01x$.
 - Non-zero gradient when the input is negative




Pool layer – Max pooling

- Reduce the spatial dimension of an image




<http://cs231n.github.io/convolutional-networks/>




Norm layer


- 4D-array



- 3D-array of 1D-vector




- Normalize each element of this 1D-vector





Norm layer

- Normalize each element of this 1D-vector





$$a_1 = \left(\left(\frac{R}{\sqrt{R^2 + G^2 + B^2}} \right), \left(\frac{G}{\sqrt{R^2 + G^2 + B^2}} \right), \left(\frac{B}{\sqrt{R^2 + G^2 + B^2}} \right) \right)$$



Local layer


- Fully connected

| | | |
|-------|-------|-------|
| x_1 | x_2 | x_3 |
| x_4 | x_5 | x_6 |
| x_7 | x_8 | x_9 |




Softmax-linear layer

- Softmax output function
- Cost measure for softmax




Softmax output function

- Soft continuous version of Max Function
- Forces $\sum(\text{Output of NN}) = 1$.





$$Y_i = \frac{e^{Z_i}}{\sum_{j \in \text{group}} e^{Z_j}}$$




Softmax output function

- $\frac{\delta Y_i}{\delta Z_i} = Y_i (1 - Y_i)$
- Nice Simple derivative
- Even though Y_i depends of Z_i ,
 - Derivative
 - for an individual neuron
 - of an O/P in respect to I/P is just $Y_i (1 - Y_i)$



Cost measure for softmax

- Cross entropy cost function
 - $\mathcal{C} = -\sum_j t_j \log Y_j$
 - Negative log probability of correct answer
 - Maximise the log probability of getting answer right
 - Very big gradient when O/P is 1 and target is 0
- $\frac{\delta \mathcal{C}}{\delta Z_i} = Y_i - T_i$
 - Slope is -1 when target values and actual value is opposite





HYPERPARAMETERS




Learning Rate

- How fast the network trains
- High learning rate
 - Convergence or global minimum finding is problem
- Low learning rate
 - High training times

Learning Rate decay


- Learning rate decay means the learning rate decreases over time
 - higher learning rate is well suited to get close to the global minimum
 - small learning rate is better at fine tuning the global minimum
- Several way
 - Exponential decay, reduction by factor of n
 - Function to decrease the learning rate by 4%

Overfitting or Underfitting



Weight Penalty



- Adding λ to penalise
 - Keeps weight small
 - Big error derivatives



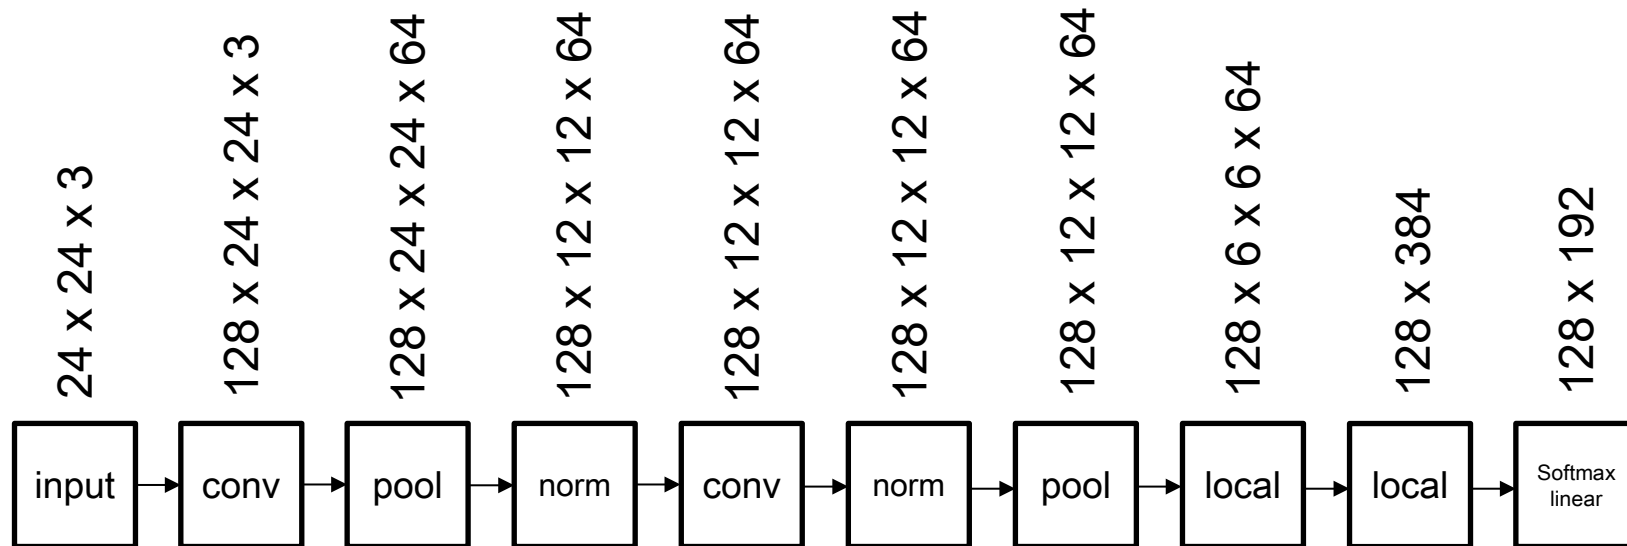
- $C = E + \frac{\lambda}{2} \sum_{i=1} w_i^2$
- $\frac{\partial C}{\partial w_i} = \frac{\partial E}{\partial w_i} + \lambda w_i$
- When $\frac{\partial C}{\partial w_i} = 0$;
 - $w_i = -\frac{1}{\lambda} \frac{\partial E}{\partial w_i}$
 - So, at minimum of cost function if $\frac{\partial E}{\partial w_i}$ is big, the weights are big

Weight Penalty - Advantages

- Preventing network from the weights it does not need
 - Don't have a lot of weights not doing anything
 - So output changes more slowly as input changes.
- Putting half the weight on each and not on one



Structure of the CNN we used



Output: 128×2



PROBLEM


The data

- Images of cats and dogs
- File format is *.jpg
- Color space is RGB



Data

- 25,000 images
 - 12,500 of dogs
 - 12,500 of cats
- Avg. file size
 - 22.34 kB



Train and test data

- Split data
 - Train data
 - 20,000 images (80 percent)
 - Divided into 5 batches containing 4,000 each
 - Test data
 - 5,000 images (20 percent)

Process images

- Resize to $32 * 32 * 3 = 3,072$
- Convert to array
 - $25,000 * 3,073$



dog1.jpg



cat10.jpg

Process images

- Resize to $32 * 32 * 3$
- Convert to array
 - $25,000 * 3,073$
- Example
 - 1; 22; 11; 123; ...
 - 0; 256; 255; 0; ...



dog1.jpg



cat10.jpg




Random distorsion






EVALUATION


Learning rate



Cross-entropy



Total loss



Total loss after 100k steps roughly above 0.1

Summarize

- Understood a CNN
- Trained it
 - Total loss is decreasing over time
 - Precision just above 0.1



QUESTIONS

Quellen

- <http://cs231n.github.io/convolutional-networks/>
- https://www.tensorflow.org/tutorials/deep_cnn/
- Maas, Andrew L., Awni Y. Hannun, and Andrew Y. Ng. "Rectifier nonlinearities improve neural network acoustic models." *Proc. ICML*. Vol. 30. No. 1. 2013.