



A Convolutional Neural Network for Image Classification of Cats and Dogs

Status update




Structure

- Neural Nets (NN)
- Convolution NN (CNN)
- Problem
- Evaluation
- Aims



NEURAL NETS

Introduction – Neural Nets



Quelle: <http://cs231n.github.io/convolutional-networks/>




CONVOLUTIONAL NN

TensorFlow




- Developed by Google Brain Team
- Use cases
 - Handwritten patterns, image recognition, Word2Vec
- Input data
 - Audio, image, text
- Used techniques
 - Linear classifiers, NN


Structure of the CNN we used



Input layer



Convolutional layer - Filter




Output

1



$$\begin{aligned}
 &1*0 + 2*0 + 0*0 + \\
 &0*0 + 0*2 + 0*1 + \\
 &0*0 + 0*2 + 0*2 + \\
 &1 = 1
 \end{aligned}$$

<http://cs231n.github.io/convolutional-networks/>




Convolutional layer - Parameters

- Input volume size
- Number of filters
- Filter size
- Zero padding





Convolutional layer – Activation function

- Sigmoid
 - Not telling in which direction should we move in.
 - Non-differentiability at certain points




$$f(x) = \frac{1}{1 + e^{-\beta x}}$$




Convolutional layer – Activation function

- Rectified linear
 - *Element Wise*: $\max(0, x)$
 - Leaky ReLU
 - If $x < 0$, Output = $0.01x$.
 - Non-zero gradient when the input is negative




Pool layer




Pool layer – Max pooling


- Reduce the spatial dimension of an image




<http://cs231n.github.io/convolutional-networks/>




Norm layer



Local layer



Softmax-linear layer





PROBLEM

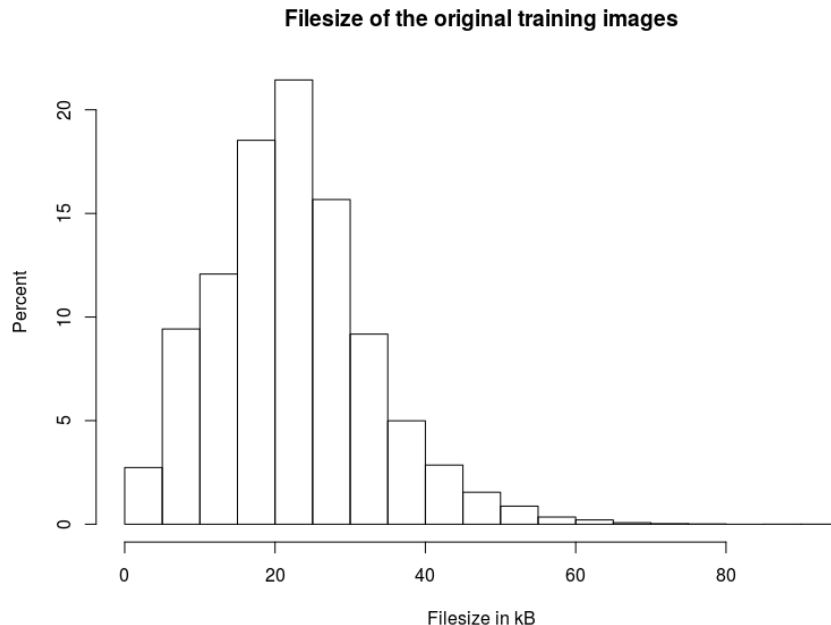
The data

- Images of cats and dogs
- File format is *.jpg
- Color space is RGB



Training data

- 25,000 images
 - 12,500 of dogs
 - 12,500 of cats
- Avg. file size
 - 22.34 kB



Test data

- 12,500 images
 - x of dogs
 - y of cats
 - $x + y = 12,500$

Process images

- Resize to $32 * 32 * 3$
- Convert to array
 - $25,000 * 3,073$



dog1.jpg



cat10.jpg

Process images

- Resize to $32 * 32 * 3$
- Convert to array
 - $25,000 * 3,073$



dog1.jpg



cat10.jpg





EVALUATION



AIMS



Aims

- Removing normalization layer



QUESTIONS

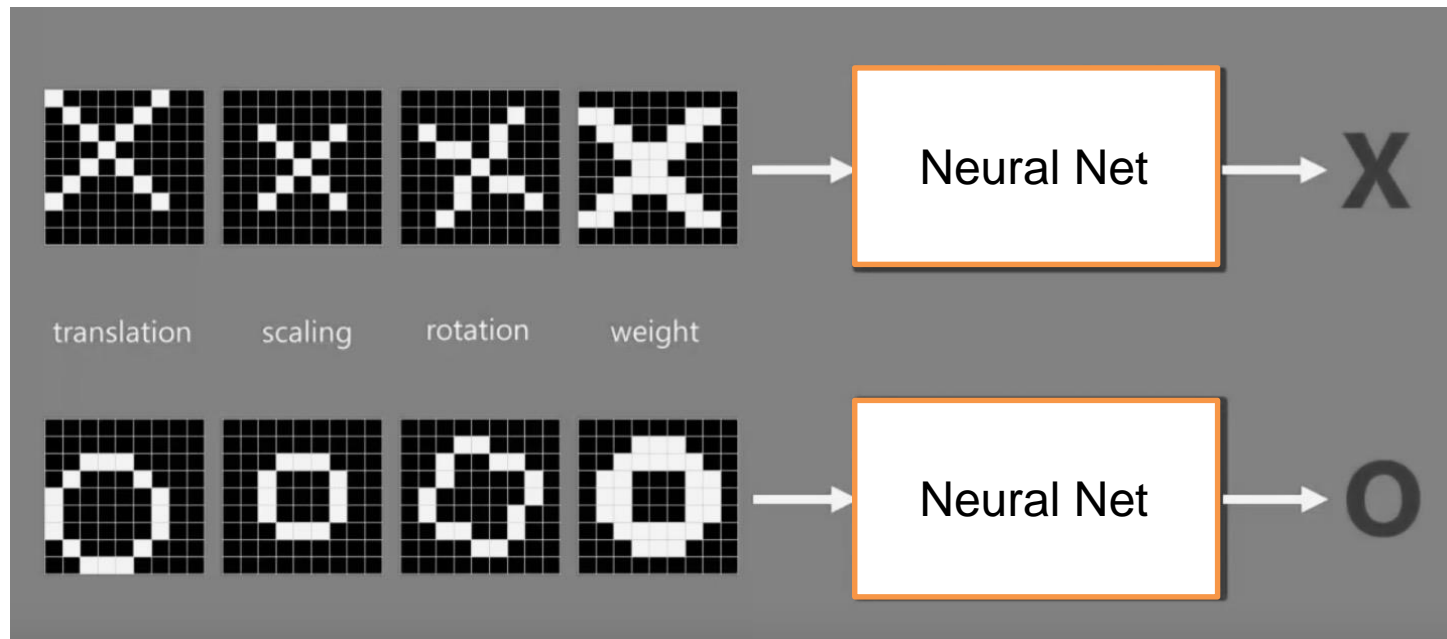
Quellen

- <http://cs231n.github.io/convolutional-networks/>
- https://www.tensorflow.org/tutorials/deep_cnn/
- Maas, Andrew L., Awni Y. Hannun, and Andrew Y. Ng. "Rectifier nonlinearities improve neural network acoustic models." *Proc. ICML*. Vol. 30. No. 1. 2013.


Today's Talk

- Problem Statement
- Introduction to Deep Learning
 - Layers In Deep Learning
- TensorFlow
 - Methods
 - DataSets
 - Training Time
- TensorFlow (TF)
 - Data-Structure for TensorFlow
- Implementation in TF
 - Inputs
 - Prediction
 - Training
 - Evaluation
- Results
- Future Works

Problem Statement - Explained



Introduction – Neural Nets



Quelle: <http://cs231n.github.io/convolutional-networks/>

Mathematical view


- Input, Weights
 - Compute Sigmoid
 - Measure how much we missed called Err
 - Multiply Err by the Sigmoid slope
 - Update Weights
- $l_0 = X_i, W = \text{rand}()$
 - $l_1 = \text{Sig}(X_i \cdot W)$
 - $\text{Err} = l_0 - l_1$
 - $\Delta l_1 = \text{Err} \cdot \Delta(\text{Sig}(\text{Err}))$
 - $W = W + (l_0 \cdot \Delta l_1)$

Gradient Descent

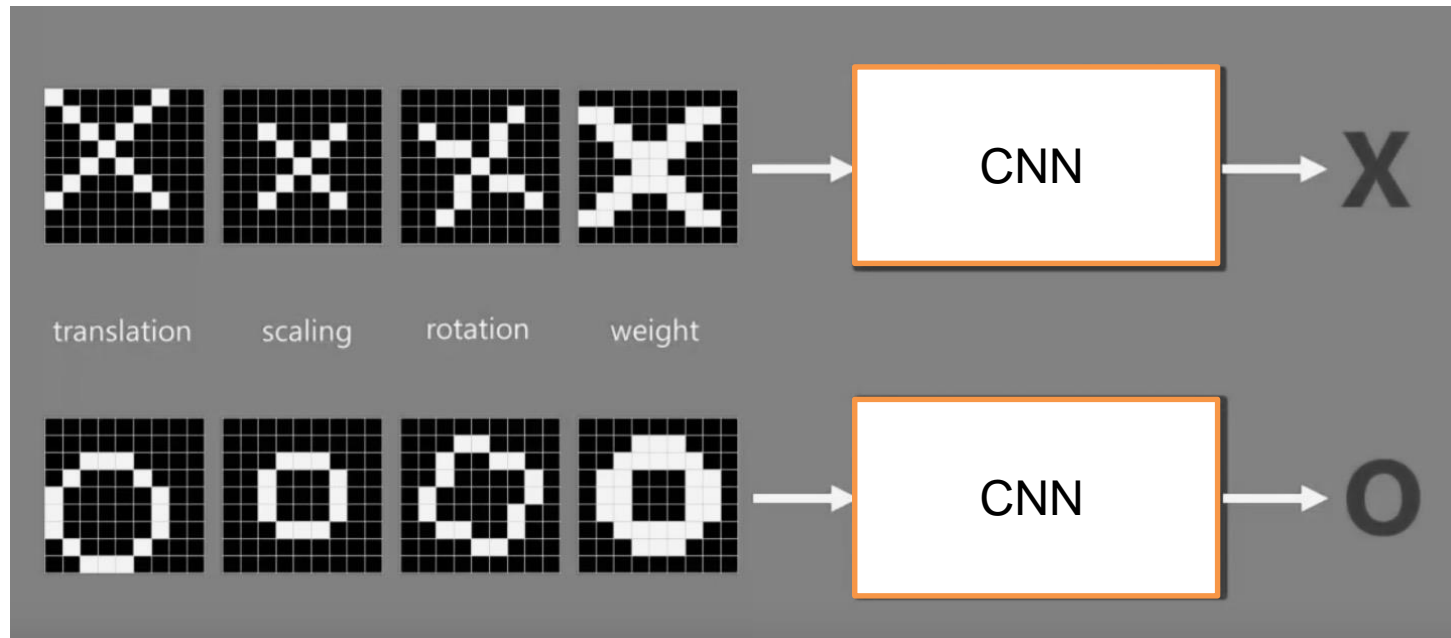
- Pushing down weights to push errors down
 - Move weights in negative direction of gradient
- $\Delta w_i = h(y - y')x_i$. *Perceptron Case*
 - *No Thresholding, finite convergence but in linear cases*
- $\Delta w_i = h(y - a)x_i$ *Activation Case*
 - *Thresholding, more robust to non-linear cases*
 - *Converge to a limit only to a local optimum*

Why not just Neural Nets?

- Input $32*32*3 = 3072$
- Weights $3072*N$
- Biases N
- So,
 - Full connectivity is wasteful
 - Huge number of parameters
 - Loss of spatial information
 - How 3072 input signals represent $32*32*3$ matrix ?
 - Deconvolution




Convolutional Neural Network (CNN)





CNN Layers

- INPUT
- CONV
- RELU
- POOL
- FC




Output




1

$$\begin{aligned}
 &1*0 + 2*0 + 0*0 + \\
 &0*0 + 0*2 + 0*1 + \\
 &0*0 + 0*2 + 0*2 + \\
 &1 = 1
 \end{aligned}$$




Result= Stack of Filtered images



Input Image



Features



Stack of Images


Convolutional Neural Network

Convolutional Filter Size

- Uneven dimensions such as 3×3 , 5×5 ...
 - To reduce spatial dimension
 - Padding can undo dimensionality reduction
- Number of conv layers:
 - More conv layers with small filters
 - This makes the decision function more discriminative


CNN – POOL

- reduce the spatial dimension of an image




Activation functions

- Why Sigmoid?
 - Not telling in which direction should we move in.
 - Non-differentiability at certain points



Sigmoid

$$f(x) = \frac{1}{1 + e^{-\beta x}}$$



CNN – RELU

- *Element Wise*: $\max(0, x)$
- Leaky ReLU by by Maas et al.
 - if $x < 0$, Output = $0.01x$.
 - non-zero gradient when the input is negative

CNN Parameters

- Can be trained on an endless amount of parameters:
 - learning rate, learning rate decay
 - momentum
 - filter size, number of convolutional layers
 - activation functions (relu, leakyrelu)
 - weight decay and dropouts

Learning Rate

- how fast the network trains
- High learning rate
 - Convergence or global minimum finding is problem
- Low learning rate
 - High training times

Learning Rate decay

- Learning rate decay means the learning rate decreases over time
 - higher learning rate is well suited to get close to the global minimum
 - small learning rate is better at fine tuning the global minimum
- Several ways to do it:
 - Exponential decay, reduction by factor of n
 - GoogLeNet: function to decrease the learning rate by 4%

Momentum

- Rolling ball gains speed downwards the hill
- So, velocity to the learning rate in a given direction
 - With consistent gradient
- Convolutional neural networks commonly use a momentum value of 0.9

Batch Normalization (BN)

- *batch normalization* of the input to the activation function of each neuron
- normalizing the training batch after certain layers
 - Reducing amount of retraining
 - input to the activation function across each training batch has a mean of 0 and a variance of 1.
- Example
 - $BN \text{ of } \sigma(Wx + b) = \sigma(BN(Wx + b))$ Where $BN = \frac{X_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$


Is BN enough ?

- No:
 - Activation function limited to a prescribed normal distribution
- Adding γ and β -> Learnable parameters
 - γ , undoes the batch normalizing transform
 - β , a new shift parameter

$$BN = \gamma \left(\frac{X_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \right) + \beta$$


Overfitting vs Underfitting

Example: Logistic regression




$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

(g = sigmoid function)




$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$$



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$$

Weight Penalty



- Adding extra term to cost function to penalise
 - Keeps weight small
 - Big error derivatives



- $C = E + \frac{\lambda}{2} \sum_i w_i^2$
- $\frac{\partial C}{\partial w_i} = \frac{\partial E}{\partial w_i} + \lambda w_i$
- When $\frac{\partial C}{\partial w_i} = 0$;
 - $w_i = -\frac{1}{\lambda} \frac{\partial E}{\partial w_i}$
 - So, at minimum of Cost function if $\frac{\partial E}{\partial w_i}$ is big, the weights are big

Weight Penalty - Advantages

- Preventing network from the weights it does not need
 - Don't have a lot of weights not doing anything
 - So output changes more slowly as input changes.
- Putting half the weight on each and not on one



Cifar10 Weight Decay

- `weight_decay = tf.mul(tf.nn.l2_loss(var), wd, name='weight_loss')`
- `tf.add_to_collection('losses', weight_decay)`



TU Clausthal

Weight Penal

TensorFlow



- Developed by Google Brain Team
- Use cases
 - Handwritten patterns, image recognition, Word2Vec
- Input data
 - Audio, image, text
- Used techniques
 - Linear classifiers, NN

Limitations

- 150 images viewed
 - 1 duplicate
 - 1 wrong content
- Training lasts long
- Reduced image size

Activation Functions

- Non-Linear
 - sigmoid, tanh, elu, softplus, and softsign
- continuous but not everywhere differentiable functions
 - relu, relu6, crelu and relu_x

TF Implementation: Prediction

■ Cifar10.Inference()

- Conv1: convolution and rectified linear activation.
- Pool1: max pooling.
- Norm1: local response normalization.
- Local3: fully connected layer with rectified linear activation.
- Local4: fully connected layer with rectified linear activation.
- Softmax_linear: Linear transformation to produce logits.



Cost Functions

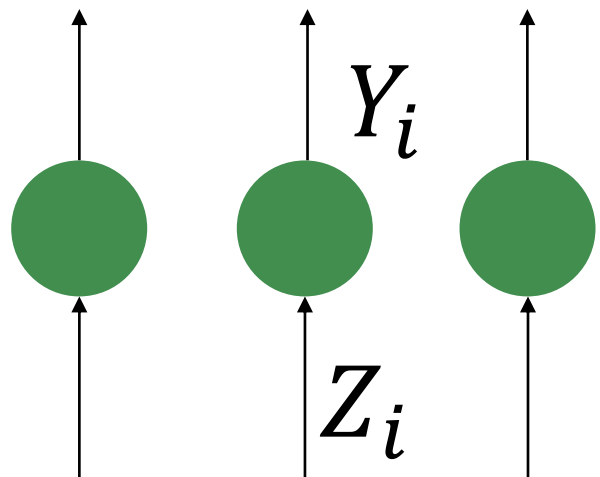
- Squared Error Measure
- Softmax Cross-entropy Function

Squared Error Measure Function

- $Error = \frac{1}{2} (Y_{actual} - Y_{predicted})^2$
- Drawbacks
 - No gradient to get from 0.000...1 to 1.
 - To do so it will take quite longer.
 - Deprives NN of probability information.

Softmax Output Function

- Soft continuous version of Max Function
- Forces $\sum(\text{Output of NN}) = 1$.


$$Y_i = \frac{e^{Z_i}}{\sum_{j \in group} e^{Z_j}}$$

Derivative Softmax

- $\frac{\delta Y_i}{\delta Z_i} = Y_i (1 - Y_i)$
- Nice Simple derivative
- Even though Y_i depends of Z_i ,
 - Derivative
 - for an individual neuron
 - of an I/P in respect to O/P is just $Y_i (1 - Y_i)$

Cost Measure for Softmax Output

Function

$$C = - \sum_j t_j \log Y_j$$

- Negative log probability of correct answer
- Maximise the log probability of getting answer right

Advantages to Squared Error Measure

- $C = -\sum_j t_j \log Y_j$
- Very big gradient when:
 - Target value is 1.
 - Actual output is 0.
- Balance between
 - Steepness of $\frac{dC}{dy}$ and
 - Flatness of $\frac{dy}{dZ}$

$$\frac{\partial C}{\partial Z_j} = \sum_j \frac{\partial C}{\partial y_j} \frac{\partial y_j}{\partial Z_j}$$



Hyperparameters

- Learning Rate





Results


- Precision 0.83



CNN – Conv Layer



Feature 1



1/9

Problem Statement

I am **DOG**
...No No...
I am **CAT**



Ratio

